

## 1. What Agents-for-python actually gives you

From the repo + docs, the **Python SDK** is basically:

- An **Activity model & types** ([microsoft-agents-activity](#)): strongly typed activities/messages, replacing the old Bot Framework schema. ([GitHub](#))
- **Hosting primitives** ([microsoft-agents-hosting-core](#), [microsoft-agents-hosting-aiohttp](#)):
  - `AgentApplication[TurnState]` and `TurnContext`
  - HTTP hosting via aiohttp and `start_agent_process` / `CloudAdapter` to receive `/api/messages` from channels. ([Microsoft Learn](#))
- **Channel integrations** ([microsoft-agents-hosting-teams](#)): glue to host the same agent in **Teams** / **M365 Copilot** / **Webchat** and other channels via the shared Activity protocol. ([GitHub](#))
- **State & storage**: `MemoryStorage`, plus optional Azure Blob and CosmosDB storage packages for durable conversation state. ([GitHub](#))
- **Auth**: `microsoft-agents-authentication-msal` for MSAL-based auth, plus the **Agents Playground** test tool for local debugging. ([GitHub](#))
- **Copilot Studio client** ([microsoft-agents-copilotstudio-client](#)): easy way to call Copilot Studio agents from your Python agent or vice versa. ([GitHub](#))

A minimal Python agent looks roughly like this (from the quickstart):

```
from microsoft_agents.hosting.core import (
    AgentApplication,
    TurnState,
    TurnContext,
    MemoryStorage,
)
from microsoft_agents.hosting.aiohttp import CloudAdapter
from start_server import start_server

AGENT_APP = AgentApplication[TurnState](storage=MemoryStorage(), adapter=CloudAdapter())

async def _help(context: TurnContext, _: TurnState):
    await context.send_activity(
        "Welcome to the Echo Agent sample. Type /help or send a message."
    )

AGENT_APP.conversation_update("membersAdded")(_help)
AGENT_APP.message("/help")(_help)

@AGENT_APP.activity("message")
```

```

async def on_message(context: TurnContext, _):
    await context.send_activity(f"you said: {context.activity.text}")

if __name__ == "__main__":
    start_server(AGENT_APP, None)

```

This gives you a working chat agent reachable via `/api/messages`, which you can attach to Teams / M365 / the playground. (Microsoft Learn)

**What it explicitly does *not* give you:**

- No built-in **LLM orchestration / tool calling / planning** (that lives in Azure AI Agent Service, Semantic Kernel, or your own code). (GitHub)
- No **data-science primitives** (datasets, summarization, forecasting, plotting, backtesting).
- No out-of-the-box **multi-step workflow engine** or task graphs.
- No opinionated **validation / Pydantic pipeline** like your in-house DSPy agent.

So: think of Agents-for-python as:

A modern Bot Framework for M365/Copilot/Teams with better hosting & storage, into which you can plug *any* LLM / agent stack (Azure AI Agent Service, Semantic Kernel, DSPy, Claude API, etc).

For your forecasting copilot, this would primarily replace “**Gradio + FastAPI**” as the chat front-door; the inner data-science agent would still look very much like what you’ve already designed.

---

## 2. “Approach matrix” view: how it scores on your 13 challenges

Using your rating scheme (– + ) and the 13 challenges from *agentic-workflow-challenges.md*, here’s an “Agents-for-python” column analogous to the existing matrix.

**Assumption:** We’re talking about an **in-house forecasting agent built on top of Agents-for-python**, with similar effort to your DSPy prototype: the SDK handles channels/state; your code handles LLM + tools.

**Summary table**

#	Challenge	M365 Agents SDK (Python)
1	Faithful Execution	Code-first model; you can enforce strong contracts & deterministic execution, but validation pipeline is yours to build.
2	Dataset Context & Summarization	No dataset abstractions; you must implement registry + summaries like today. SDK is mostly neutral here.
3	Scope Control (stay in “forecasting”)	Activity routing + explicit handlers make it easy to constrain scope; LLM still needs instructions but plumbing helps.
4	Result Presentation	Good: Teams/Webchat/Adaptive Cards for tables + text; charts/images possible but require custom front-end work.
5	Ambiguity Handling & Defaults	Depends entirely on your LLM/orchestrator; SDK doesn't help or hurt, it just delivers activities to your code.
6	Session Data Management & Caching	Strong: built-in storage providers + <code>TurnState</code> let you persist per-conversation state cleanly.
7	History / Context Management	You get stored state and conversation IDs, but you still need to decide which bits go into LLM prompts.
8	Error Communication & Recovery	SDK surfaces errors cleanly; turning them into user-friendly explanations / recovery flows is up to your logic.
9	Multi-Step Planning	No planning primitives; you must implement task graphs or use Azure AI Agent Service/Semantic Kernel explicitly.
10	Data Quality & Validation	Nothing built-in for datasets; same situation as a plain FastAPI backend—you'd port your current checks.
11	Responsiveness & Progress Feedback	Async I/O, typing indicators, and channel features let you show progress, but you must wire progress messages yourself.
12	Tool Design & Granularity	Full freedom: no 10-action limit, no MCP schema constraints; you design functions/sub-agents as you like.
13	LLM Configuration & Control	You own the LLM calls (Azure OpenAI, others) so you control model, temperature, retries, routing, A/B tests, etc.

### How that compares conceptually to your other approaches

Using your existing matrix as reference:

- Versus **ChatGPT GPTs**:
  - You **lose** the zero-code convenience and built-in Code Interpreter plots, but
  - You **gain** full control over tools, LLM config, and channel UX; no 10-action limit, no vendor-defined prompting.
- Versus **Claude MCP**:
  - Claude brings “free” multi-step planning and ambiguity handling; Agents-for-python doesn’t.
  - But with Agents-for-python you get **first-class Teams/M365 integration** and a more traditional, testable Python service instead of a desktop-only tool. (GitHub)
- Versus **In-house DSPy**:
  - Capability-wise, they’re quite similar: both are “you write the agent logic” frameworks.
  - DSPy helps with **prompt programming, structured outputs, and declarative agents**; the Microsoft SDK helps with **channels, hosting, and enterprise plumbing**.

So if you literally plug your existing DSPy agent (or a refactored version of it) behind a Microsoft AgentApplication, many of your 13-challenge scores would be **identical to DSPy**, because the heavy lifting is done by your logic, not by the SDK.

Where the **SDK itself** really moves the needle relative to your current architecture:

- **Session & state** (Challenge 6): instead of hand-rolled `agent_state` in Gradio, you get a formal state API, storage providers, and per-conversation TurnState. (Microsoft Learn)
  - **Result presentation** (Challenge 4): you can leverage Teams cards, M365 Copilot surfaces, and custom webchat front-ends instead of just Gradio.
  - **Multi-channel deployment & trust** (orthogonal to the 13 challenges but important): you can reach users via Teams, M365 Copilot, and webchat, reusing the same agent container. (GitHub)
- 

### 3. “Equivalent approach” architecture for your forecasting copilot

To match your existing architecture docs (DSPy, Claude MCP, ChatGPT GPTs), here’s how a “**M365 Agents SDK (Python) + forecasting backend**” approach would look.

### 3.1 High-level design

Teams / M365 Copilot / Webchat

Microsoft Agents SDK (Python)

- AgentApplication[TurnState]
- CloudAdapter + aiohttp hosting
- Storage (Memory/Blob/Cosmos)

Forecasting Orchestrator (your code)

- Existing Pythia tools (predict, backtest, explain, plot)
- Dataset registry & summarization
- Clarifier / defaults logic
- Validation + Pydantic models

Pythia API / data sources

**Key idea:** The Microsoft agent becomes your **front-door / conversation container**, and for “real work” it calls the same Python orchestration you’ve already designed (DSPy-style or not).

### 3.2 Sketch of the core interaction

#### 1. Hosting & routing (Agents-for-python):

```
# app.py
from microsoft_agents.hosting.core import (
    AgentApplication, TurnState, TurnContext, MemoryStorage
)
from microsoft_agents.hosting.aiohttp import CloudAdapter
from start_server import start_server

from forecasting_agent import handle_forecasting_turn # your module

AGENT_APP = AgentApplication[TurnState](
    storage=MemoryStorage(),
    adapter=CloudAdapter()
)

@AGENT_APP.activity("message")
async def on_message(context: TurnContext, state: TurnState):
    user_text = context.activity.text or ""
    # Delegate to your forecasting orchestrator
    reply, artifacts = await handle_forecasting_turn(
```

```

        text=user_text,
        conversation_id=context.activity.conversation.id,
        state=state,
    )
    await context.send_activity(reply)
    # Optionally send separate messages/cards for plots, tables, etc.

if __name__ == "__main__":
    start_server(AGENT_APP, auth_configuration=None)

2. Your forecasting orchestrator (very close to today):

# forecasting_agent.py (pseudo)
async def handle_forecasting_turn(text, conversation_id, state):
    """
    Wrap your existing agent stack (DSPy/LLM+tools) behind a simple function.
    """

    # 1. Recover cached dataset/predictions from TurnState or external cache
    current_ctx = state.get_value("forecasting_ctx") or {}

    # 2. Run your interpreter + clarifier + tool-selection logic
    #     (could still be DSPy; could be Azure AI Agent Service; could be Claude via API)
    new_ctx, artifacts = await run_forecasting_workflow(
        user_message=text,
        context=current_ctx,
    )

    # 3. Save updated state
    state.set_value("forecasting_ctx", new_ctx)

    # 4. Build a natural-language reply for the channel
    reply_text = build_reply_text(artifacts)

    return reply_text, artifacts

```

Here, `Agents-for-python` doesn't know about datasets, forecasting horizons, or Pythia at all; it just:

- Accepts messages from Teams/M365/webchat,
- Maintains conversation state, and
- Invokes your forecasting orchestrator to handle each turn.

This preserves almost all of the work you've already done on the **13 challenges** (structured params, dataset registry, display flags, etc.) and swaps only the **outer shell** (Gradio/CLI) for **M365 channels**.

### 3.3 Where you'd still need significant work

If you went “all in” on Agents-for-python as *the* agent framework (rather than “just a front-door” to your existing logic), you’d need to:

- **Re-implement your DSPy multi-agent logic** (Interpreter, Clarifier, Orchestrator) as either:
  - Plain Python code calling an LLM (e.g. Azure OpenAI), or
  - Azure AI Agent Service / Semantic Kernel orchestrations wired into the AgentApplication pipeline. (GitHub)
- **Bring over your Pydantic contracts & validation** for Pythia tools and datasets (Challenge 1 & 2).
- **Recreate your progress callbacks and display flags** in terms of “send one or more activities to the channel” (Challenge 4 & 11).
- **Implement planning** if you want first-class multi-step workflows (Challenge 9); Azure’s new durable agents on .NET have more guidance here, but for Python you’d largely be rolling your own. (Microsoft Learn)

So compared to your **current DSPy prototype**, an Agents-for-python approach is:

- **More work on the “agent core”** if you abandon DSPy’s abstractions.
  - **Less work on channels & enterprise plumbing**, especially if you care about Teams, M365 Copilot, or Copilot Studio integration.
- 

## 4. So... how well does it cover *your* needs?

Grounding in your documents and challenges:

- If your **goal** is “**Forecasting copilot in Teams / M365 Copilot / webchat**”, → **Agents-for-python is a strong front-door**, and you can reuse almost all of your existing forecasting logic behind it.
- If your **goal** is “**minimize new agent code**”, → It doesn’t buy you much on the 13 hard problems (faithful execution, dataset context, planning, data quality, etc.). Those still need the same custom logic you’ve already designed in DSPy / Claude / GPTs.
- If your **goal** is “**replace DSPy/Claude/GPTs with a single Microsoft stack that handles everything end-to-end**”, → Agents-for-python by itself is **not** that stack; you’d also lean on Azure AI Agent Service / Semantic Kernel / durable agents, which are currently more mature on .NET than Python. (Microsoft Learn)

My practical take:

- Treat **Agents-for-python as an *integration layer***:

- Keep your current in-house agent design (or Claude/GPT orchestration) for the 13 technical challenges.
  - Wrap it in a Python AgentApplication so it can run in Teams / Copilot / Webchat with proper auth, storage, and observability.
- In your **approach comparison**, it probably deserves its own row, but conceptually it's closer to “**In-house DSPy fully developed**” than to Claude MCP or ChatGPT GPTs:
  - Full control, high flexibility, higher engineering cost, low vendor lock-in on the *LLM*, but **strong lock-in on the channel stack** (M365).